



The Abdus Salam
International Centre
for Theoretical Physics



Shared Memory Programming Paradigm

Ivan Girotto – igirotto@ictp.it

Information & Communication Technology Section (ICTS)
International Centre for Theoretical Physics (ICTP)



OUTLINE


- The Shared Memory Programming Paradigm
- Processes and Threads
- The OpenMP Standard
- Parallel Programming with OpenMP
- Hands-on

The Recommended Approach

- Parallelism is no longer an option for only either larger scale problems or improve the time of response
- It is inescapable to exploit current & next generations of compute processors




Multi-core system Vs Serial Programming




Xeon E5650
hex-core
processors
(12GB - RAM)

Multi-core system Vs // Programming



Xeon E5650
hex-core
processors
(12GB - RAM)

Multi-CPUs & Multi-cores system





Xeon E5650
hex-core
processors

Main Memory

Dual Socket (Westmere) - 24GB RAM

Processes and Threads



STACK

SHARED MEMORY

TEXT

compiled code (a.out)

System
env argv argc
auto variables for main()
auto variables for func()
<i>available for stack growth</i>
malloc.o (lib*.so)
printf.o (lib*.so)
<i>available for heap growth</i>
Heap (malloc arena)
global variables
"...%d..."
malloc.o (lib*.a)
printf.o (lib*.a)
file.o
main.o func(72,73)
crt0.o (startup routine)

High memory

mfp – frame pointer (for main)

stack pointer
(grows downward if func()
calls another function)

brk point

uninitialized data (bss)

initialized data

library functions if
statically linked
(not usual case)

ra (return address)

Low memory

Stack illustrated after the call
func(72,73) called from main(),
assuming func defined by:
func(int x, int y) {
 int a;
 int b[3];
 /* no other auto variables */
}

Assumes int = long = char * of
size 4 and assumes stack at high
address and descending down.

Expanded view of the stack

Offset from current
frame pointer (for
func())

frame pointer
points here

EBP

stack pointer
(top of stack)
points here

ESP

Stack



main() auto variables	Contents
+12	y
+8	x
+4	return address
0	caller's frame pointer
-4	a
-8	b[2]
-12	b[1]
-16	b[0]

All auto variables and parameters
are referenced via offsets from the
frame pointer.

The frame pointer and stack pointer
are in registers (for fast access).

When func returns, the return value
is stored in a register. The stack pointer
is move to the y location, the code
is jumped to the return address (ra),
and the frame pointer is set to mfp
(the stored value of the caller's frame
pointer). The caller moves the return
value to the right place.

Processes and Threads



Multi-threading - Recap

- A thread is a (**lightweight**) process - an instance of a program plus its own data (private memory)
- Each thread can follow its own flow of control through a program
- Threads can share data with other threads, but also have private data
- Threads communicate with each other via the shared data.
- A *master thread* is responsible for co-ordinating the threads group



The Abdus Salam
International Centre
for Theoretical Physics



OpenMP™


OpenMP (*Open spec. for Multi Processing*)

- OpenMP **is not a computer language**
 - Rather it works in conjunction with existing languages such as standard Fortran or C/C++
- Application Programming Interface (API)
 - that provides a **portable** model for parallel applications
 - Three main components:
 - Compiler **directives**
 - **Runtime library** routines
 - **Environment variables**

OpenMP Parallelization

- OpenMP is directive based
 - code (can) work without them
- OpenMP can be added incrementally
- OpenMP only works in shared memory
 - multi-socket nodes, multi-core processors
- OpenMP hides the calls to a threads library
 - less flexible, but much less programming
- **Caution:** write access to shared data can easily lead to race conditions and incorrect data

OpenMP Parallelization



- Thread-based Parallelism
- Explicit Parallelism
- Fork-Join Model
- Compiler Directive Based
- Dynamic Threads

Getting Started with OpenMP

- OpenMP's constructs fall into 5 categories:
 - Parallel Regions
 - Work sharing
 - Data Environment (scope)
 - Synchronization
 - Runtime functions/environment variables
- OpenMP is essentially the same for both Fortran and C/C++

Directives Format

- A directive is a special line of source code with meaning only to certain compilers.
- A directive is distinguished by a sentinel at the start of the line.
- OpenMP sentinels are:
 - Fortran: **!\$OMP** (or **C\$OMP** or ***\$OMP**)
 - C/C++: **#pragma omp**


OpenMP: Parallel Regions

- For example, to create a 4-thread parallel region:
 - each thread calls `foo(ID,A)` for **ID = 0 to 3**

Each thread redundantly executes the code within the structured block

thread-safe routine: A routine that performs the intended function even when executed concurrently (by more than one thread)

```
double A[1000];
omp_set_num_threads(4);
#pragma omp parallel
{
    int ID =omp_get_thread_num();
    foo(ID,A);
}
printf( "All Done\n" );
```



How many threads?

The number of threads in a parallel region is determined by the following factors:

- Use of the `omp_set_num_threads()` library function
- Setting of the `OMP_NUM_THREADS` environment variable
- The implementation **default**

Threads are numbered from 0 (master thread) to N-1.

OpenMP runtime library

`OMP_GET_NUM_THREADS()` – returns the current # of threads.


`OMP_GET_THREAD_NUM()` - returns the id of this thread.


`OMP_SET_NUM_THREADS(n)` – set the desired # of threads.

`OMP_IN_PARALLEL()` – returns .true. if inside parallel region.

`OMP_GET_MAX_THREADS()` - returns the # of possible threads.

Memory footprint





Simple C OpenMP Program

```
#include <omp.h>
```

```
#include <stdio.h>
```

```
int main ( ) {
```

```
    printf("Starting off in the sequential world.\n");
```

```
    #pragma omp parallel
```

```
    {
```

```
        printf("Hello from thread number %d\n", omp_get_thread_num() );
```

```
    }
```

```
    printf("Back to the sequential world.\n");
```

```
    return 0;
```

```
}
```

Variable Scooping

- All existing variable still exist inside a parallel region
 - by default SHARED between all threads
- But work sharing requires private variables
 - PRIVATE clause to OMP PARALLEL directive
 - Index variable of a worksharing loop
 - All declared local variable within a parallel region
 - The FIRSTPRIVATE clause would initialize the private instances with the contents of the shared instance
- Be aware of the sharing nature of static variables

Exploiting Loop Level Parallelism

Loop level Parallelism: parallelize only loops

- Easy to implement
- Highly readable code
- Less than optimal performance (sometimes)
- Most often used

Parallel Loop Directives

- Fortran do loop directive
 - `!$omp do`
- C\C++ for loop directive
 - `#pragma omp for`
- These directives do not create a team of threads but assume there has already been a team forked.
- If not inside a parallel region shortcuts can be used.
 - `!$omp parallel do`
 - `#pragma omp parallel for`

Parallel Loop Directives continued

- These are equivalent to a parallel construct followed immediately by a worksharing construct.

!\$omp parallel do

Same as

!\$omp parallel

...

!\$omp do

#pragma omp parallel for

Same as

#pragma omp parallel

...

#pragma omp for

How is OpenMP Typically Used?

- OpenMP is usually used to parallelize loops:

Split-up this loop between multiple threads


```
void main()
{
    double Res[1000];

    for(int i=0;i<1000;i++) {
        do_huge_comp(Res[i]);
    }
}
```

Sequential program

```
void main()
{
    double Res[1000];
    #pragma omp parallel for
    for(int i=0;i<1000;i++) {
        do_huge_comp(Res[i]);
    }
}
```

Parallel program



Work-Sharing Constructs

- Divides the execution of the enclosed code region among the members of the team that encounter it.
- Work-sharing constructs **do not** launch new threads.
- No implied barrier upon entry to a work sharing construct.
- However, there is an implied barrier at the end of the work sharing construct (unless **nowait** is used).

Work Sharing Constructs - example

Sequential code

```
for(i=0;I<N;i++) { a[i] = a[i] + b[i];}
```


OpenMP // Region

```
#pragma omp parallel
{
    int id, i, Nthrds, istart, iend;
    id = omp_get_thread_num();
    Nthrds = omp_get_num_threads();
    istart = id * N / Nthrds;
    iend = (id+1) * N / Nthrds;
    for(i=istart;I<iend;i++) {a[i]=a[i]+b[i];}
}
```

OpenMP Parallel
Region and a work-
sharing for construct

```
#pragma omp parallel
#pragma omp for schedule(static)
for(i=0;I<N;i++) { a[i]=a[i]+b[i];}
```

schedule(static [,chunk])




- Iterations are divided evenly among threads
- If chunk is specified, divides the work into chunk sized parcels
- If there are N threads, each thread does every Nth chunk of work.

```
!$OMP PARALLEL DO &
!$OMP SCHEDULE(STATIC,3)
```

```
DO J = 1, 36
Work (j)
END DO
```

```
!$OMP END DO
```

schedule(dynamic [,chunk])



- Divides the workload into chunk sized parcels.
- As a thread finishes one chunk, it grabs the next available chunk.
- Default value for chunk is one.
- More overhead, but potentially better load balancing.

```
!$OMP PARALLEL DO &
!$OMPSCHEDULE(DYNAMIC,1)
```

```
DO J = 1, 36
Work (j)
END DO
```

```
!$OMP END DO
```

The Schedule Clause SCHEDULE (type [,chunk])

- The schedule clause effects how loop iterations are mapped onto threads

`schedule(static [,chunk])`

- Deal-out blocks of iterations of size “chunk” to each thread

`schedule(dynamic [,chunk])`

- Each thread grabs “chunk” iterations off a queue until all iterations have been handled

`schedule(guided [,chunk])`

- Threads dynamically grab blocks of iterations. The size of the block starts large and shrinks down to size “chunk” as the calculation proceeds

`schedule(runtime)`


- Schedule and chunk size taken from the OMP_SCHEDULE environment variable

Compute PI

$$\int_0^1 \frac{1}{1+x^2} dx = \arctan(x) \Big|_0^1 = \arctan(1) - \arctan(0) = \arctan(1) = \frac{\pi}{4}$$

$$\pi = 4 \int_0^1 \frac{1}{1+x^2} dx$$

Integrate, i.e determine area under function numerically using slices of $h * f(x)$ at midpoints



No Wait Clauses

- No wait: if specified then threads do not synchronise at the end of the parallel loop.
- For Fortran, the END DO directive is optional with NO WAIT being the default.
- Note that the nowait clause is incompatible with a simple parallel region meaning that using the composite directives will not allow you to use the nowait clause.

OpenMP: Reduction(**op** : **list**)

- The variables in “list” must be shared in the enclosing parallel region.
- Inside a parallel or a worksharing construct:
 - A local copy of each list variable is made and initialized depending on the “op” (e.g. 0 for “+”)
 - pair wise “op” is updated on the local value
 - Local copies are reduced into a single global copy at the end of the construct.

OpenMP: A Reduction Example

```
#include <omp.h>
#define NUM_THREADS 2
void main ()
{
    int i;
    double ZZ, func(), sum=0.0;

    omp_set_num_threads(NUM_THREADS);
    #pragma omp parallel for reduction(+:sum) private(ZZ)
    for (i=0; i< 1000; i++){
        ZZ = func(i);
        sum = sum + ZZ;
    }
}
```

if CLAUSE

We can make the parallel region directive itself conditional.

Fortran: *IF (scalar logical expression)*

C/C++: *if (scalar expression)*

```
#pragma omp parallel if (tasks > 1000)
{
    while(tasks > 0) donexttask();
}
```



The Abdus Salam
**International Centre
for Theoretical Physics**



SYNCHRONIZATION

OpenMP: How do Threads Interact?


- OpenMP is a shared memory model.
 - Threads communicate by **sharing variables**.
- Unintended sharing of data can lead to **race conditions**:
 - race condition: when the program's outcome changes as the threads are scheduled differently.
- To control race conditions:
 - Use **synchronization** to protect data conflicts.
- Synchronization is expensive so:
 - Change how data is stored to minimize the need for synchronization.

Note that updates to shared variables:

(e.g. $a = a + 1$)

are *not* **atomic**!

If two threads try to do this at the same time, one of the updates may get overwritten.



Barrier

Fortran	- !\$OMP BARRIER
C\C++	- #pragma omp barrier

- This directive synchronises the threads in a team by causing them to wait until all of the other threads have reached this point in the code.
- Implicit barriers exist after work sharing constructs. The `nowait` clause can be used to prevent this behaviour.

Critical

- Only one thread at a time can enter a **critical** section.

Example: pushing and popping a task stack

```
!$OMP PARALLEL SHARED(STACK),PRIVATE(INEXT,INEW)
```

```
...
```

```
!$OMP CRITICAL (STACKPROT)
```

```
    inext = getnext(stack)
```

```
!$OMP END CRITICAL (STACKPROT)
```

```
    call work(inext,inew)
```

```
!$OMP CRITICAL (STACKPROT)
```

```
    if (inew .gt. 0) call putnew(inew,stack)
```

```
!$OMP END CRITICAL (STACKPROT)
```

```
...
```

```
!$OMP END PARALLEL
```

Atomic

- **Atomic** is a special case of a critical section that can be used for certain simple statements

Fortran: **!\$OMP ATOMIC**
statement

where *statement* must have one of these forms:

$x = x \text{ op } \text{expr}$, $x = \text{expr op } x$, $x = \text{intr} (x, \text{expr})$ or
 $x = \text{intr}(\text{expr}, x)$

op is one of +, *, -, /, .and., .or., .eqv., or .neqv.

intr is one of **MAX**, **MIN**, **IAND**, **IOR** or **IEOR**



Non Parallelizzabile

- Show an example of Instruction dependency



The Abdus Salam
International Centre
for Theoretical Physics



HANDS-ON ON THREADED LIBRARIES



- Source:
<http://www.inl.gov/computing/tutorials/openMP/#ProgrammingModel>



1) Create the submission script

```
[hpc01@java2 ~]$ more sub_script.sh
```

```
#!/bin/bash
#PBS -q mhpc
#PBS -l nodes=1:ppn=16
#PBS -l walltime=2:00:00

cd $PBS_O_WORKDIR

export OMP_NUM_THREADS=8

#for Intel Compiler

module load intel/2013

./my_prog.x
```

3 fundamental steps
for job execution

2) Submit the script

```
[hpc01@java2 ~]$ qsub sub_script.sh
70.java2.grid.lipi.go.id
```

3) Monitor job execution

```
[hpc01@java2 ~]$ qstat
```

Job id	Name	User	Time Use	S	Queue
70.java2	test_pbs	ictp	0 R		nogpu